

Backend Coding Challenge  
Jan 2017

### **Programming Exercise**

You have been asked to build a server to reassemble UDP packets into N discrete messages. Design an in-memory data model (eg. No database or filesystem). Leverage the data model to store the fragments from which the completed message can be reassembled. Provide an implementation of the consumer/reassembly process. The output of the system should be the completed message number and its corresponding sha256 hash. If all fragments arrived, the hash should line up with the sha256 hash output by the emitter. If the process crashes, there is no expectation of recovery for data in flight.

Fragments delivered in the packet should be stored in the data model. A packet with the EOF flag can arrive before other packets in the same message, and all packets need to be included. If the completed message has not been received within 30 seconds, an error should be issued. The error should indicate which pieces of the specific message did not arrive. (eg. perhaps byte ranges 3-5 and 72-79 did not arrive).

The implementation may be written in any language using corresponding tech stacks. If no extra credit is attempted, one instantiation of your process must run and consume data off the UDP port.

Please include a readme describing the design of the implementation, how to execute the server, and any known deficiencies or limitations. If a partial solution is delivered, describe which requirements are not met and how you would evolve the code to meet those requirements.

### **Multi-threading extra credit (Medium Difficulty)**

Due to UDPs lack of reliable delivery and the high rate of incoming packets, you will need to run 4 consumer threads concurrently. This will keep incoming packets from being lost while the server is processing the last packet. The data model must be enhanced to support parallel worker threads. Each consumer will be listening for UDP communications over a single, shared port.

The implementation may be written in any language using corresponding tech stacks, however, it is a requirement that 4 threads of your process run concurrently consuming data off the UDP port.

### **Fault-tolerant and multi-process extra credit (Difficult)**

Your implementation should run in four discrete processes. Your implementation should also be fault-tolerant and limit data loss in the event of a crash or a restart. This involves persisting in-flight data to disk or a database.

### **What is included in this ZIP**

Included is a node.js client program that will emit 10 messages of variable size on UDP port 6789 on the local loopback. The sha256 for each message is computed and output as follows:

```
> node udp_emitter.js  
Emitting message #1 of size:1022990
```



sha256:ffc846e16558fa7501f6deb71ee9c4e1d9db0eeb5421deadafbcd9e605f249e8  
Message #3 length: 865334  
sha256:7b60a643bbcecc48f5277f0c58006f0644e557e1fa25cee220142d04b18f0402  
Message #4 length: 961815  
sha256:c39d75497a7d9c7d758fec1bd98edf318f4819b06f185c80ab2b22b044b40b6c  
Message #5 length: 470530  
sha256:2c349e023f6536e874f2695b32c562680cd488520ea902b4134e4e417756b347  
Message #6 length: 260786  
sha256:98f7b66ba315e40f81a98c8072336ec360441c939f7f3328293ac677406486b2  
Message #7 Hole at: 66  
Message #7 Hole at: 2324  
Message #7 Hole at: 3722  
Message #7 Hole at: 4416  
Message #7 Hole at: 5490  
Message #7 Hole at: 9873  
Message #7 Hole at: 12500  
Message #7 Hole at: 16474  
Message #8 length: 459077  
sha256:d3644218d5d96c41f5cf18843adf71eb7dbbb4a9cdb2af1b310a8d8c53763f25  
Message #9 length: 22273  
sha256:fd428f2c7e7bec03811d595320blad69ee79e143710e8ae330616321bc64a4bf  
Message #10 length: 10409  
sha256:30a0022f6e89c4d28691ce15387201ffbbcc99f30773584d390c5984330557ef